



# Data Engineering

## LTAT.02.007

Ass Prof. Riccardo Tommasini

Fabiano Spiga, Hassan Eldeeb, Mohamed Ragab

## Graph Databases (Neo4j)


lab 05

# Agenda

- Relational vs. graph databases
  - o Which to use and when?
- From Relational to Graph
- What is Neo4j?
- The Property Graph Model (PGM)
- Cypher Query Language
- Cypher In Action
  - o MovieGraph Dataset in neo4j
- Cypher Query Clauses and Examples



# Relational vs. graph databases



VS.



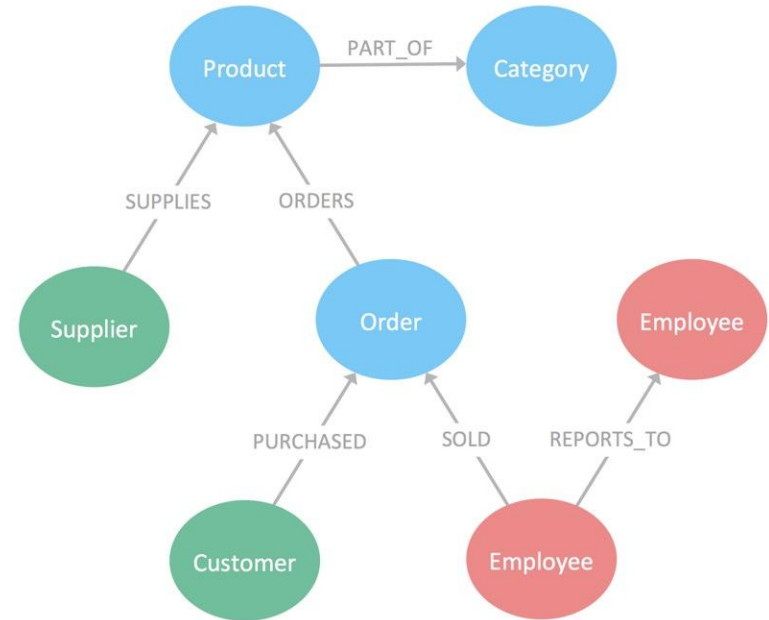
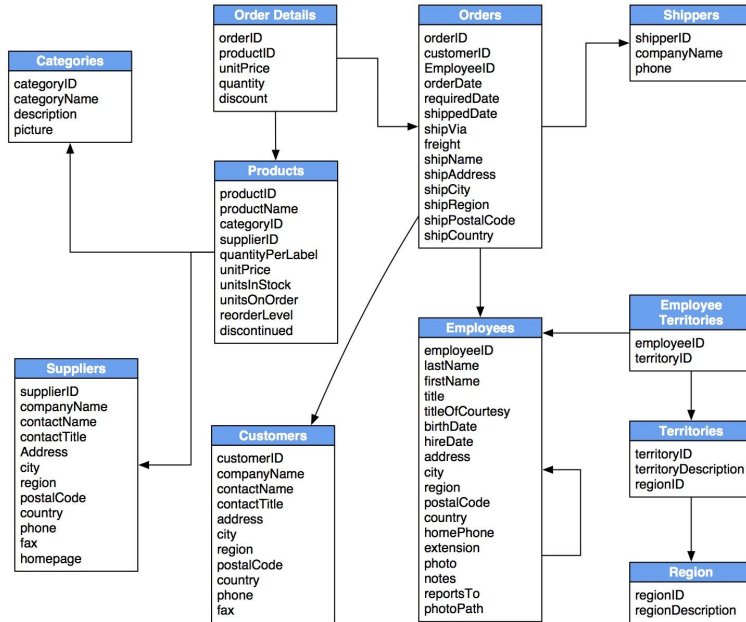
## Good For:

- Well-understood data structures that don't change too frequently
- Known problems involving discrete parts of the data, or minimal connectivity

## Good For:

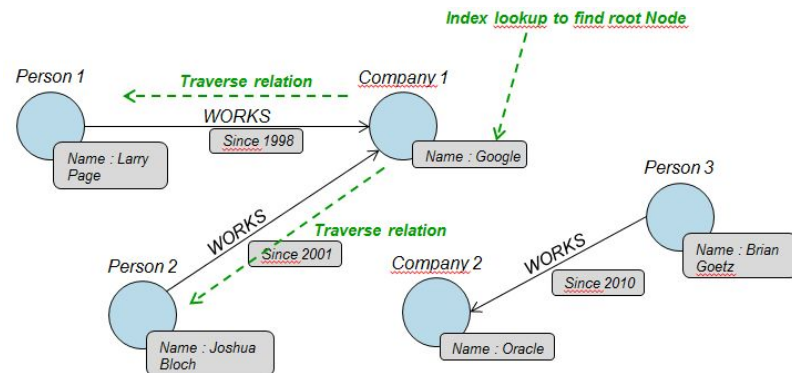
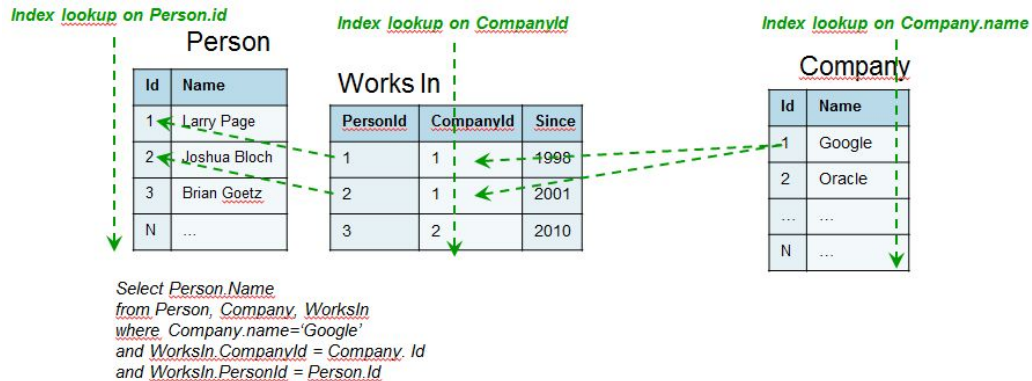
- Dynamic systems where the data topology is difficult to predict
- Dynamic requirements that evolve with the business
- Problems where the relationships in data contribute meaning and value

# From Relational to Graph




<https://neo4j.com/business-edge/relational-to-relationships/>

# From Relational to Graph



# Graph Databases



neo4j

OrientDB



ArangoDB



Amazon Neptune



TITAN



AllegroGraph

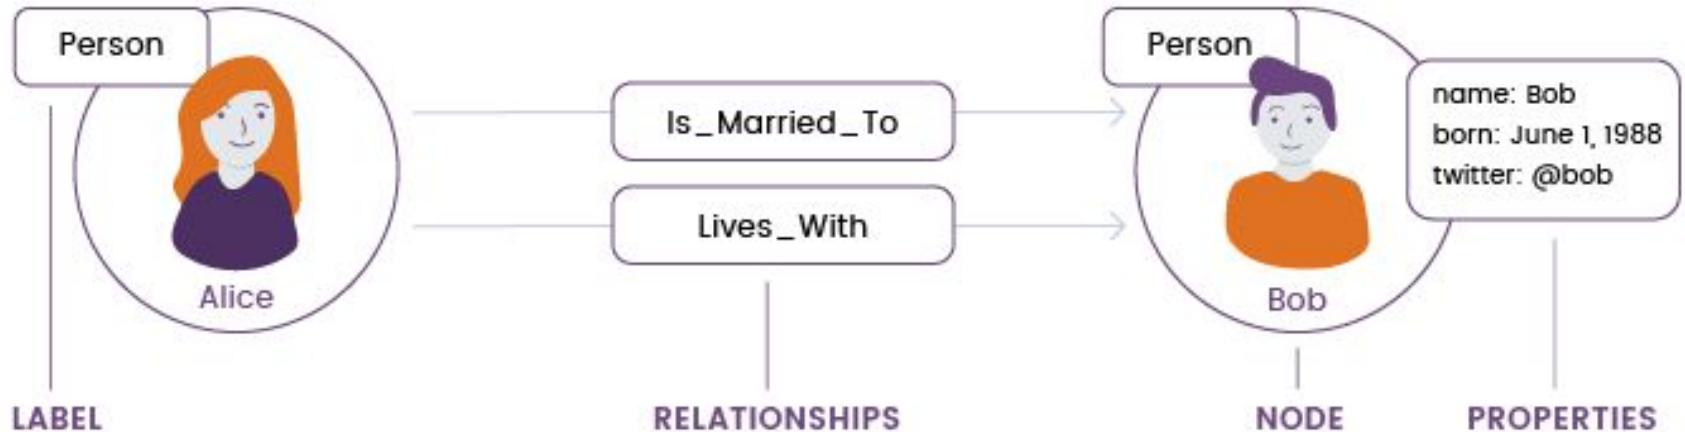
Franz Inc.

# Neo4j



- A graph **database management system** developed by Neo4j, Inc.
- Supports **ACID** (Atomicity, Consistency, Isolation, Durability) properties.
- It implements a **Property Graph Model** to store the data.
- Its implementation in **Java** also makes it widely usable, you still can access it using **other languages** using available **drivers**.
- It uses **Cypher** as a declarative query language.


# Property Graph Model







# What is Cypher?

- Graph Query Language for Neo4j, and other GDBs.
- Aims to make graph querying simple.
- Cypher is **Declarative QL**
  - specify starting point specify desired outcome
  - algorithm adaptable based on query
- Based on **graph Pattern matching** (**ASCII-art patterns**).
- Cypher is Familiar for **SQL users**.





# Pattern Matching

- Queries in form of graph pattern





# Pattern Matching

- Select at least Start Node





# Pattern Matching

- Our first Match





# Pattern Matching

- This is not match!



# Pattern Matching

- This is also non-Match.
- Why?!



# ASCII-art patterns

Directed relationship



$(A) \dashrightarrow (B)$

# ASCII-art patterns

Undirected relationship




(A)--(B)



# ASCII-art patterns


Specified relationship



$(A) -[:Loves]->(B)$

# ASCII-art patterns


Joined Paths



(A) --> (B) --> (C)

# ASCII-art patterns

multiple Paths



>> (A) -->(B) --> (c), (A) --> (c)

>> (A) -->(B) --> (c) <-- (A)

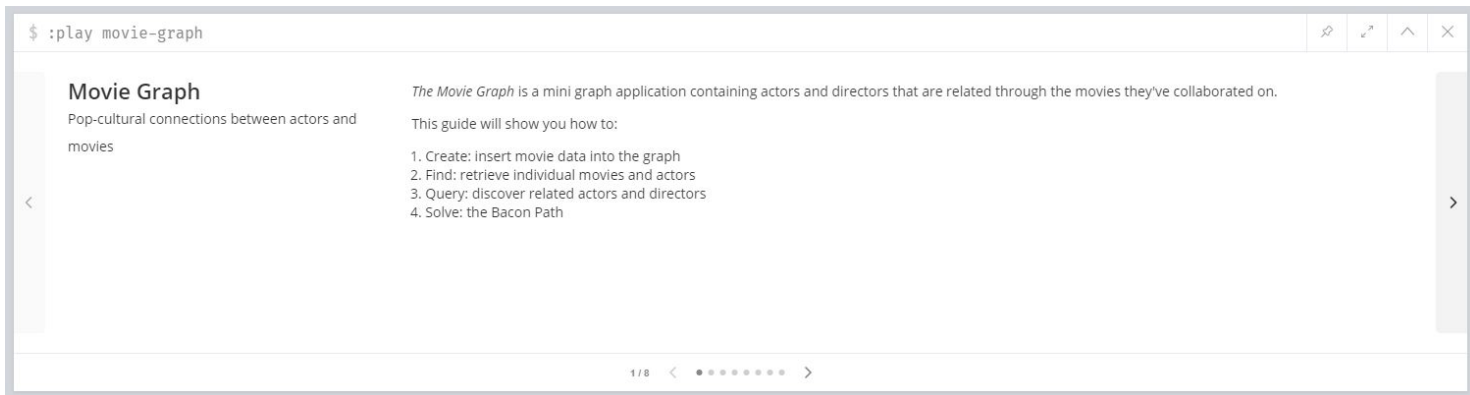
# Cypher In Action



<https://s3.amazonaws.com/images.seroundtable.com/google-hands-dirty-1436184854.jpg>

# Movie Dataset Example

- Excerpt of actors, directors, producers,, etc and related movies.
- Available on your installed Neo4j versions.
- In your neo4j browser run:  
:play movie-graph



# Loading the movie dataset

\$ :play movie-graph

## The Movie Graph

### Create

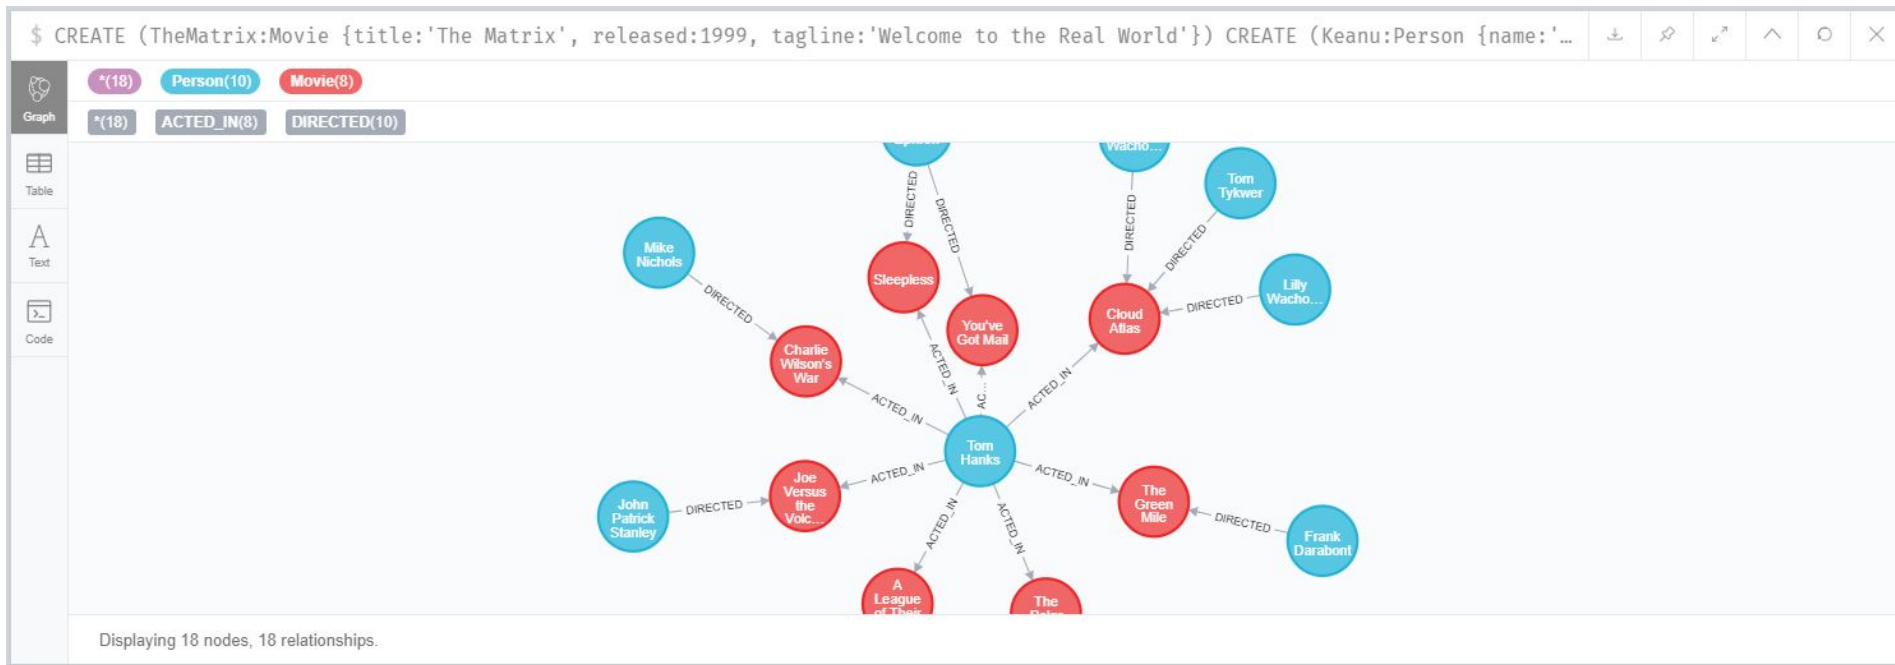
To the right is a giant code block containing a single Cypher query statement composed of multiple CREATE clauses. This will create the movie graph.

1. Click on the code block
2. Notice it gets copied to the editor above ↑
3. Click the editor's play button to execute
4. Wait for the query to finish  
WARNING: This adds data to the current database, each time it is run!

:help ☒ cypher ☐ CREATE

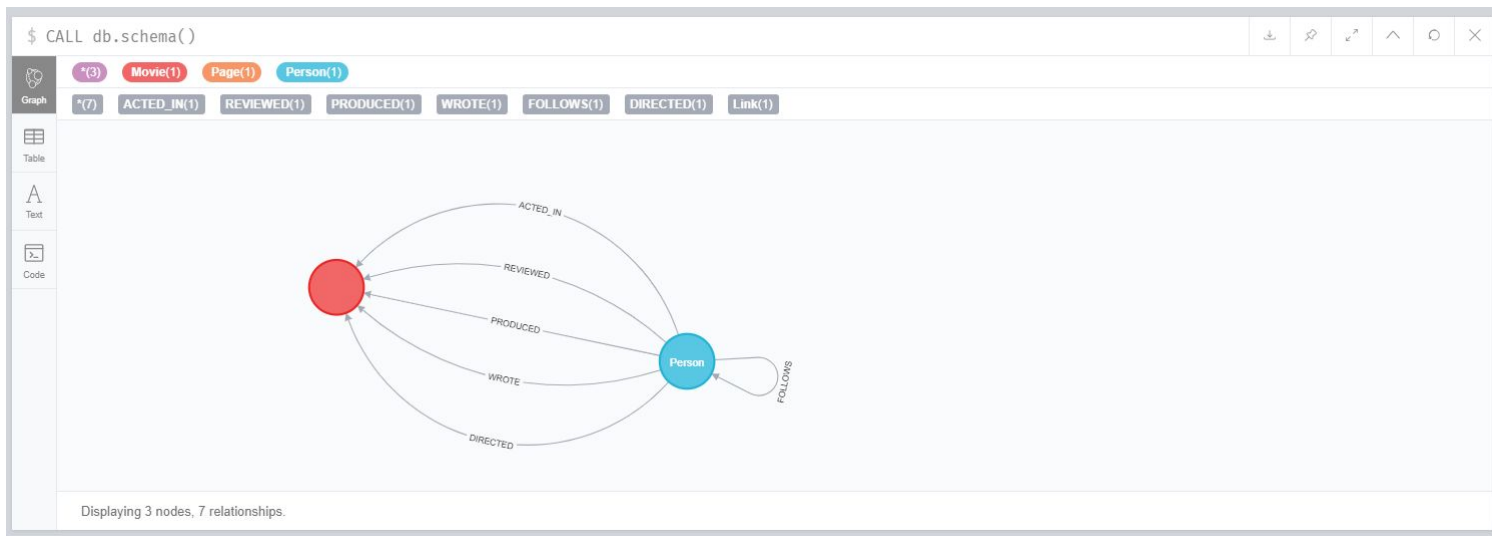
```
CREATE (TheMatrix:Movie {title:'The Matrix', released:1999, tagline:'Welcome to the Real World'})
CREATE (Keanu:Person {name:'Keanu Reeves', born:1964})
CREATE (Carrie:Person {name:'Carrie-Anne Moss', born:1967})
CREATE (Laurence:Person {name:'Laurence Fishburne', born:1961})
CREATE (Hugo:Person {name:'Hugo Weaving', born:1960})
CREATE (LillyW:Person {name:'Lilly Wachowski', born:1967})
CREATE (LanaW:Person {name:'Lana Wachowski', born:1965})
CREATE (JoelS:Person {name:'Joel Silver', born:1952})
CREATE
(Keanu)-[:ACTED_IN {roles:['Neo']}]>(TheMatrix),
(Carrie)-[:ACTED_IN {roles:['Trinity']}]>(TheMatrix),
(Laurence)-[:ACTED_IN {roles:['Morpheus']}]>(TheMatrix),
(Hugo)-[:ACTED_IN {roles:['Agent Smith']}]>(TheMatrix),
(LillyW)-[:ACTED_IN {roles:['The Architect']}]>(TheMatrix),
(LanaW)-[:ACTED_IN {roles:['The Architect']}]>(TheMatrix),
(JoelS)-[:ACTED_IN {roles:['The Architect']}]>(TheMatrix)
```

# Results of data loading



# Show the DB Schema

- From the menu Favorites>Common Procedures choose> show meta-graph.
- Or directly run the query: `CALL db.schema()`





# CREATE Nodes

- This clause is used to create nodes, relationships, and properties.
- **Creating a Single node**
  - CREATE (node\_name)
- **Creating Multiple Nodes**
  - CREATE (node1) , (node2)
- **Creating a Node with a Label**
  - CREATE (node:label)
  - CREATE (Messi:player)
- **Creating a Node with Multiple Labels**
  - CREATE (node:label\_1:label\_2:.. . . . label\_n)
  - CREATE (Messi:person:player)
- **Create Node with Properties**
  - CREATE (node:label { key1: value, key2: value, . . . . . })
  - CREATE (Messi:player{name: "Lionel Messi", YOB: 1987, POB: "Rosario"})

# CREATE Relationships

- **Creating Relationships**

- CREATE (node1)-[:RelationshipType]->(node2)
- CREATE (Messi:player{name: "Lionel Messi", YOB: 1987, POB: "Rosario"})
- CREATE (Arg:Country {name: "Argentina"})
- CREATE (Messi)-[r:BORN\_IN]->(Arg)

- **Creating a Relationship with Label and Properties**

- CREATE (node1)-[label:Rel\_Type {key1:value1, key2:value2, . . . n}]-> (node2)

- **Creating a Complete Path**

- CREATE p = (Node1 {properties})-[:Relationship\_Type]-> (Node2 {properties})[:Relationship\_Type]->(Node3 {properties}) RETURN p

# MATCH

- The MATCH clause introduces the pattern we are looking for.
- Node patterns can contain labels and properties.

```
MATCH (e:Employee) RETURN e.name, e.surname
```

- Finding nodes by relationships

```
MATCH (e:Employee)-->(cc:CostCenter) RETURN *
```

- \* Returns all named nodes, relationships and identifiers

# MERGE

- MERGE command is a **combination** of CREATE command and MATCH command.
- It searches for a given pattern in the graph. If it **exists**, then it returns the results.
- If it does **NOT exist** in the graph, then it creates a new node/relationship and returns the results.
- Syntax: MERGE (node: label {properties . . . . . })
  - MERGE (node:label) RETURN node
- OnCreate and OnMatch
  - MERGE (node:label {properties . . . . . })
  - ON CREATE SET property.isCreated ="true"
  - ON MATCH SET property.isFound ="true"

# Filtering properties

- we can get all the employees that have a relation with a specific cost center, for example `CC1`.
- Properties are expressed within curly brackets.

```
MATCH (e:Employee) --> (:CostCenter {code: 'CC1' })  
RETURN e
```

- Here, as you can notice, we omitted the `cc` variable (we don't need it!).

# Matching by the Relationship

- Relationship expressions must be specified in square brackets.
- To filter the employees who **belong to** a specific cost center, we have to specify the relationship type:

```
MATCH (n) -[:BELONGS_TO]-> (:CostCenter { code: 'CC1' } )  
RETURN n
```

# OPTIONAL MATCH

- Search for the pattern described in it, while using nulls for missing parts of the pattern.
- OPTIONAL MATCH is similar to the match clause, the only difference being it returns null as a result of the missing parts of the pattern.

```
MATCH (a:Movie { title: 'Cloud Atlas' })  
OPTIONAL MATCH (a)-->(x)  
RETURN x
```

- Returns **null**, since the node has no outgoing relationships.

# WHERE

- Use a predicate to filter. Note that WHERE is always part of a MATCH, OPTIONAL MATCH or WITH clause.
- Putting it after a different clause in a query will alter what it does.

```
WHERE n.property <> $value
```

```
MATCH (e:Employee) --> (cc:CostCenter)
WHERE cc.code='CC1'
RETURN e
```



# WHERE

- Filter on node label

**Match (n)**

**WHERE n:Swedish**

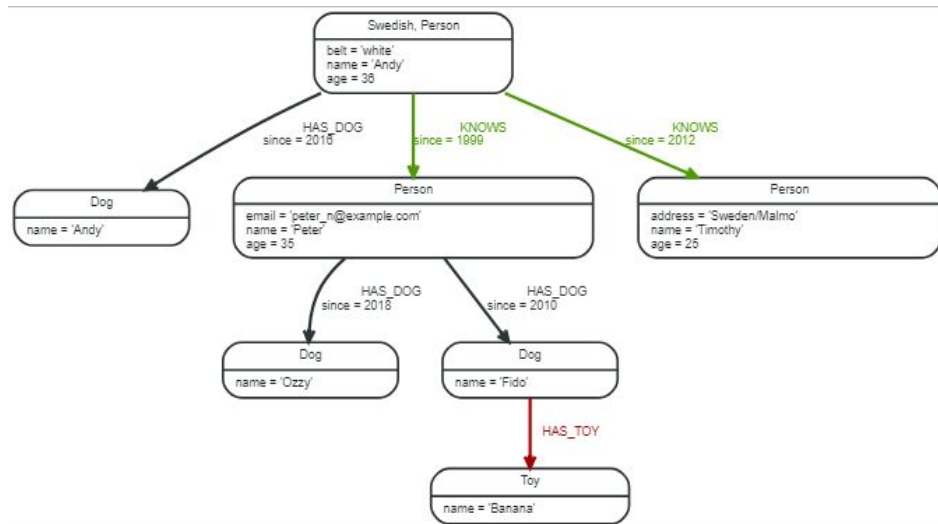
**RETURN n.name, n.age**

- Filter on a relationship property

**MATCH (n:Person) - [k:KNOWS] -> (f)**

**WHERE k.since < 2000**

**RETURN f.name, f.age, f.email**



# DISTINCT (Unique results)

- DISTINCT retrieves only unique rows depending on the columns that have been selected to output.

```
RETURN distinct expr
```

- Example:

```
MATCH (a { name: 'A' }) --> (b)
```

```
RETURN DISTINCT b
```


- The node named "B" is returned by the query, but only once.

# Paging results – LIMIT and SKIP

- Paging is necessary to avoid a lot of data from being loaded all together in a single query.

```
MATCH (b:Book)
WHERE b.title =~ '(?i).*drama.*'
RETURN b
LIMIT 20
```

```
MATCH (b:Book)
WHERE b.title =~ '(?i).*drama.*'
RETURN b
SKIP 20
LIMIT 20
```



# Sorting (Order By)

- Sorting with Cypher is exactly the same as sorting with SQL using (ORDER BY).

```
MATCH (b:Book) RETURN b.title
ORDER BY b.title
LIMIT 5
```

- A descending sort: By default sort is ascending, you can reverse the order using **DESC** clause.

```
MATCH (b:Book)
RETURN b.title
ORDER BY b.title DESC
```

# COUNT

- `count(*)`, counts the number of matching rows.
- `count(variable)`, counts the number of non-null values.
- `count(DISTINCT variable)`, count can take the DISTINCT operator, which removes duplicates from the values.

```
MATCH (b) <- [r:Vote] - (u:User)
      RETURN COUNT(*) as votes
```

# Collecting values in an array

- You can collect all the values in an array so that you can easily process them with your preferred algorithm.
- For example, the following query returns all the score votes received for books:

```
MATCH (b:Book) <-[r:Votes]- (:User)
RETURN b.title, COLLECT(r.score)
```

```
+-----+
| b.title           | COLLECT(r.score) |
+-----+
| "Epic of Gilgamesh" | [5,4,3,4,1]      |
| "The Divine Comedy" | [4,3,5,3,4]      |
+-----+
```

# WITH (Separating query parts)

- The WITH syntax is similar to RETURN.
- It separates query parts explicitly, allowing you to declare which variables to carry over to the next part.
- One common usage of WITH is to **limit** the number of entries that are then **passed on** to other MATCH clauses.

```
MATCH (user)-[:FRIEND]-(friend)
WHERE user.name = $name
WITH user, count(friend) AS friends
WHERE friends > 10
RETURN user
```

# Cypher Docs. & Cheat Sheet

- <https://neo4j.com/docs/cypher-manual/3.5/introduction/>
- <https://neo4j.com/docs/cypher-refcard/current/>
- <https://neo4j.com/docs/cypher-manual/current/>



Funny Pictures on [www.LeFunny.net](http://www.LeFunny.net)



Now, It's time to say ...

THANK YOU

**SEE YOU  
NEXT TIME!**